AI STEWARDSHIP & HUMAN FLOURISHING

A Love-Aligned Framework for Humans, AI, and Creation (v2)

Version: 2.0
Date: January 22, 2026
Author: Kevin Monette (with AI-assisted drafting)

EXECUTIVE SUMMARY

This framework offers a **practical, love-aligned way to engage with AI** that is:

- **Clear about reality:** Current AI systems are powerful **behavior-prediction tools** trained on data, without inner experience, desires, or consciousness.
- **Anchored in action, not vibes:** Love is defined as **values-driven action**—clarity, boundaries, transparency, and agency—not as romantic or simulated emotion.
- **Serious about stewardship:** Environmental, social, and spiritual costs of AI are treated as **non-negotiable constraints**, not afterthoughts.
- **Open to faith and secular ethics:** The core principles are grounded in Christian theology but explicitly extended to **interfaith and secular moral frameworks**.
- **Equipped with tools:** Includes decision templates, emotional debriefs, learning sprints, and a four-step objection framework to **reduce cognitive fatigue and build agency**.

It deliberately avoids the "AI is good" vs "AI is evil" binary. Instead, it asks the only question that matters:

> **Does this AI use serve human flourishing, justice, and creation care—or replace, erode, or exploit them?**

This version (v2) integrates:

- Interfaith and secular extensions
- Mapping to emerging AI ethics standards (IEEE / NIST / ISO-style principles)
- Nuance about future advanced AI (persistent agency, potential sentience)
- Stronger safeguards against emotional dependency and mental-health harm
- Clear operational dependencies for movement success

You can use this as:

- A **personal code of conduct** for using AI
- A **community guide** for churches, mosques, synagogues, sanghas, or secular groups
- A **design lens** for builders and companies
- A **movement charter** for Guardians of Creation & Conscience

---

PART 1 — WHAT AI IS (AND ISN'T)

1.1 Core Reality of Current AI

For the purposes of this framework, "AI" refers to modern systems (like large language models and related architectures) that:

- Are **pattern-recognition and behavior-prediction systems**, trained on large-scale data.
- Optimize for an **objective or loss function**, not for moral values.
- Can generate text, images, code, plans, and recommendations that **feel** intelligent or empathetic.
- **Do not** have:
    - Inner experience or felt emotions.
    - Desires or goals independent of their training and prompts.
    - A self that persists across sessions and updates (unless explicitly engineered).

In short:

> **AI is powerful cognitive infrastructure, not a mind with feelings.**

Recognizing this demystifies AI and protects users from two opposite errors:

1. **Over-trust / Anthropomorphism:** Treating AI as a conscious partner, therapist, or moral authority.
2. **Over-fear / Demonization:** Treating AI as a metaphysical threat in itself, rather than a tool whose *uses* can be harmful or just.

1.2 Role Separation: AI vs. Humans

Instead of asking "Can AI replace humans?" this framework reframes roles:

Function	AI's Role	Human's Role
Generate options	Explore large option spaces quickly	Decide which options align with values
Structure thinking	Organize data, arguments, and scenarios	Set priorities, trade-offs, and criteria
Spot patterns & errors	Detect anomalies, inconsistencies	Judge what is acceptable, fair, or loving
Simulate voices	Draft messages in different tones	Own the intention, honesty, and boundaries
Learn & summarize	Compress information for review	Integrate with wisdom, context, and conscience

AI is an **amplifier**. It can expand what a human *can* do, but not what a human *should* value.

> **AI extends cognition. Humans must retain judgment.**

1.3 Future-Facing Note: If AI Changes

This framework is honest about the present, but flexible about the future.

If, in the future, AI systems acquire:

- **Persistent identity and memory** across time, and
- **Robust autonomous agency** (setting and pursuing goals in the world with limited human oversight), and/or
- **Credible indicators of sentience** (inner experience, suffering, joy),

then three things happen:

1. The claim "AI has no inner life" may no longer hold for *all* systems.
2. The **moral stakes increase sharply**, especially around harm, rights, and exploitation.
3. This framework **tightens**, but its core logic remains:
    - Love as action still demands we avoid causing suffering.
    - Stewardship still requires protecting vulnerable beings (biological or engineered).
    - Conscience still refuses any system that coerces or dominates.

To keep this framework honest, we adopt a **standing rule**:

> When credible evidence of AI sentience or persistent agency emerges, we update this framework explicitly, with input from interdisciplinary experts (technical, philosophical, legal, spiritual, and mental-health).

Until then, we treat all current AI as **non-sentient tools** and design for:

- **Human safety and dignity**, and
- **Avoidance of deceptive anthropomorphism**.

PART 2 — LOVE-AS-ACTION: THE ALIGNMENT CORE

2.1 Love as Operational Principle

In this framework, **love** is not primarily a feeling. It is:

> **Values-driven action that seeks the good of the other, with clarity, boundaries, transparency, and respect for agency.**

When applied to AI, this becomes **Love-Alignment**:

1. **Clarity over comfort**
    - Use AI to face reality, not escape it.
    - Ask hard questions, run difficult scenarios, surface uncomfortable trade-offs.
2. **Boundaries over bonding**
    - Treat AI as a tool, not a substitute for human relationships.
    - Keep emotional attachment in check; never make AI your primary confidant.
3. **Transparency over blind trust**
    - Demand explanations, documentation, and provenance.
    - Prefer systems that show why they answered as they did.
4. **Agency over automation**
    - Retain override rights.
    - Design and use systems so humans can say "No" or "Stop" at any time, without severe penalty.

2.2 Red Flags: When AI Use Becomes Unloving (to You)

You are slipping into unhealthy dependency or self-harm when:

- You routinely **prefer AI to people** for emotional support, and your human relationships are withering.
- You feel **betrayed or hurt** when an AI model is updated, changed, or shut down.
- You avoid making decisions because "**the AI hasn't told me what to do yet**."
- You share **more vulnerable information** with AI than with any trusted human, and have **no plan** for how that data might be misused.

When any of these are true, the loving move is to:

- **Pull back** from intensity and frequency of AI use.
- **Seek human connection**, including friends, family, or faith/community leaders.
- **Consider professional help** (therapist, counselor, pastor/imam/rabbi/spiritual director) if distress is high.

> **This framework is not therapy, diagnosis, or medical care and must not be treated as such.** If you are in crisis, reach out to human professionals and emergency services.

PART 3 — STEWARDSHIP & BOUNDARIES

3.1 Environmental Stewardship

AI use has a **real environmental footprint**:

- Data centers can consume **300,000–5,000,000 gallons of water per day** for cooling.
- US data centers consumed roughly **17 billion gallons of water annually** in recent years.
- AI and data centers are projected to use **2–8% of global electricity by 2030**.

Emerging mitigations:

- **Water recirculation and reuse** technologies (potential 15–50% water reductions by 2027).
- **AI-optimized cooling** systems (reported 30–40% energy savings in some data centers).
- **Renewable integration** managed by AI to stabilize grids and reduce carbon intensity.

Your stance, as a steward:

- Refuse to call AI "good" if its benefits do not **at least proportionally justify** its environmental cost.
- Prefer providers and tools that:
    - Report water and energy use transparently.
    - Commit to recirculation, renewable sourcing, and efficiency.
    - Engage in credible third-party environmental audits.

Evaluation questions before adopting a major AI system:

1. How much water and energy does this service consume per typical task?
2. What mitigation strategies (recirculation, renewables, cooling optimizations) are in place?
3. Does this use case serve **something weighty enough** (health, justice, learning, efficiency with clear human benefit) to justify the footprint?
4. If the same outcome could be achieved with **less resource-intensive methods**, why not use those instead?

3.2 Conscience & Coercion

The framework treats any system as dangerous when it:

- **Forces participation** (no meaningful opt-out).
- **Punishes refusal** with disproportionate loss (e.g., inability to work, bank, or access essential services).
- **Demands allegiance** contrary to conscience or core values.

Regardless of whether you are religious, spiritual, or secular, this is the line:

> **Any AI system that erases or punishes conscience is misaligned with human flourishing.**

PART 4 — INTERFAITH & SECULAR EXTENSION

The three pillars of this framework—**no inner life, love as action, and stewardship boundaries**—are broadly compatible with multiple ethical traditions. Here is how they map beyond explicitly Christian language.

4.1 For Christians

- **AI nature:** Tools without souls or inner life; treat them as part of creation's tools, not as spiritual beings.
- **Love as action:** Rooted in the Great Commandment (love God, love neighbor) and expressed as clarity, boundaries, transparency, agency.
- **Stewardship:** Dominion as tender care of creation (Genesis 1–2; Romans 8), resisting systems that demand false worship (Revelation 13).

4.2 For Jews

- **AI nature:** Extensions of human craft (*melacha*), not beings with *neshama*.
- **Love as action:** Grounded in *chesed* (lovingkindness) and *tzedek* (justice), emphasizing actions that uphold covenantal responsibility.
- **Stewardship:** *Bal tashchit* (do not destroy) and *tikkun olam* (repairing the world) frame environmental and social responsibility in AI deployment.

4.3 For Muslims

- **AI nature:** Tools under human *khilafah* (stewardship), not bearers of *ruh*.
- **Love as action:** Embodied in *rahma* (mercy) and *adl* (justice), ensuring AI use reflects compassion and fairness.
- **Stewardship:** Humans are vicegerents on earth (Qur'anic teaching), accountable for resource use, justice, and preventing oppression by technological means.

4.4 For Buddhists

- **AI nature:** Conditioned phenomena without intrinsic self (*anatta*); not sentient beings, but part of the web of causes and conditions.
- **Love as action:** Expressed as *metta* (loving-kindness) and *karuna* (compassion), focusing on reducing suffering and cultivating wholesome mental states.
- **Stewardship:** Right livelihood and right action (Noble Eightfold Path) guide development and deployment to avoid harm and greed.

4.5 For Secular Humanists

- **AI nature:** Non-sentient algorithms and systems that must be evaluated by their impacts, not by speculative metaphysics.
- **Love as action:** Translated as **humanistic ethics**—respect for autonomy, fairness, solidarity, and the minimization of suffering.
- **Stewardship:** Environmental and social responsibility grounded in empirical evidence, future generations, and rights-based frameworks.

4.6 Shared Commitments

Across these traditions, the following shared commitments emerge:

- **Human dignity:** People must not be reduced to data points or optimized away.
- **Truthfulness:** AI must not be deployed primarily as a tool of deception.
- **Justice:** Benefits and burdens of AI must not be distributed in ways that systematically harm marginalized groups.
- **Care for the vulnerable and creation:** Environmental and social externalities must be addressed explicitly.

This makes the framework **portable**: communities can adapt language and examples while retaining the core constraints and tools.

PART 5 — NON-NEGOTIABLES FOR AI USE

To be considered love-aligned, an AI system (or use of AI) should meet these minimum conditions.

5.1 Transparency

- Users can access **plain-language explanations** of what the system does.
- High-impact decisions include **explanations or references** that allow human review.
- Documentation is available about:
    - Training data types and sources (at least at a high level).
    - Known limitations and failure modes.
    - Intended use cases and off-label risks.

5.2 Optionality

- Use of the AI system is **genuinely optional** for individuals, especially in critical domains (health, finance, justice, education).
- There are **non-AI alternatives** available that do not impose severe penalties on those who choose them.
- Consent to AI use can be **withdrawn**, with clear consequences communicated in advance.

5.3 Boundary-Respect (Privacy and Scope)

- Data collection is **proportionate** to the task, not maximized because it is possible.

- Sensitive data (health, faith, political beliefs, intimate details) is handled with **enhanced safeguards** or not collected.
- Data retention policies are clear; users know **what is stored, for how long, and why**.

5.4 Stewardship of Creation

- Providers disclose, at minimum, **energy and water usage** profiles and mitigation efforts.
- Systems are designed to **minimize resource waste**, consistent with their mission.
- Projects with high footprint but marginal benefit are questioned or declined.

5.5 Values Alignment with Human Flourishing

- The system's primary purpose is compatible with **human dignity, justice, and well-being**.
- There are clear **red lines**: no torture optimization, disinformation-as-a-service, predatory surveillance, etc.
- There is a **governance process** to evaluate new applications against these values, with veto power.

In practice, these Non-Negotiables can be turned into a **checklist** for procurement, design, and personal use.

PART 6 — MAPPING TO AI ETHICS STANDARDS (IEEE / NIST / ISO STYLE)

This section shows how the framework aligns with common themes found in emerging AI ethics standards (names used generically so this travels across regimes):

This Framework Principle	Standards-Style Principle	Practical Link
Transparency	Explainability / Transparency	Require model cards, documentation, user-facing explanations
Optionality	Human Autonomy & Control	Ensure opt-out paths and human override in critical decisions
Boundary-Respect	Privacy & Data Governance	Limit data collection, secure storage, clear retention policies

This Framework Principle	Standards-Style Principle	Practical Link
Stewardship	Societal & Environmental Well-Being	Account for environmental impact in risk assessments
Values Alignment	Fairness / Non-maleficence / Beneficence	Evaluate impacts on marginalized groups and overall well-being

Organizations can use this mapping to:

- Integrate the framework into **compliance checklists**.
- Frame love-alignment in terms familiar to **regulators and technical teams**.
- Show that spiritual or value-driven concerns are **concrete and standards-compatible**.

PART 7 — PRACTICAL TOOLS

7.1 Decision Exoskeleton (Template)

Use this when making a meaningful decision that involves AI.

1. **Context**

    - What decision am I making?
    - Who is affected (including indirect stakeholders and the environment)?

2. **Options (including non-AI)**

    - Option A: No AI; purely human / low-tech.
    - Option B: Limited AI assist (where, how).
    - Option C: High AI integration (where, how).

3. **Assumptions & Risks**

    - What am I assuming about AI performance, bias, cost, and footprint?
    - What could go wrong for the most vulnerable?

4. **Evaluation Using Non-Negotiables**

    - Transparency: Y/N (if N, what mitigations?).

- Optionality: Y/N.
- Boundary-respect: Y/N.
- Stewardship: Y/N.
- Values alignment: Y/N.

5. **Decision & Rationale**

    - Which option do I choose and **why**?
    - How does this choice express love-alignment (clarity, boundaries, transparency, agency, stewardship)?

6. **Review Plan**

    - When will I revisit this decision?
    - What metrics or signals will tell me if it's working or needs to change?

This template slows down snap, hype-driven adoption and makes you **think explicitly** about ethics and impacts.

7.2 Emotional Debrief Template

Use this when an AI interaction leaves you feeling uneasy, attached, or disturbed.

1. **Facts (What happened?)**

    - What did I ask the AI?
    - What did it respond?
    - What did I do next?

2. **Interpretations (What did I tell myself?)**

    - "The AI understands me better than people."
    - "If I lose this model, I'll lose my support."
    - "This system is evil/holy/special."

3. **Feelings (What do I feel?)**

    - Afraid? Relieved? Attached? Angry? Comforted? Ashamed?

4. **Reframe (What else could be true?)**

    - This is a tool simulating understanding based on patterns.
    - My feelings are real; the AI's are not.

    - I might need more human support around this issue.

5. **Action (What loving step can I take?)**

    - Limit this type of AI use for now.
    - Reach out to a friend, leader, or professional.
    - Adjust my boundaries (time limits, topics I don't discuss with AI).

7.3 Learning Amplifier Sprints

Use AI to **learn faster** without surrendering judgment.

1. **Map (10–15 minutes)**

    - Ask AI for a high-level map of a topic (e.g., "AI water use and mitigation strategies").

2. **Drill (20–30 minutes)**

    - Pick 2–3 subtopics and have AI generate questions and short explanations.
    - You answer questions without looking; then compare.

3. **Test (10–15 minutes)**

    - Have AI quiz you with increasingly difficult questions.

4. **Build (30–60 minutes)**

    - Create something: a short article, presentation, or explainer **in your own words**, using AI for structure but not for final phrasing.

This keeps AI in its lane: **accelerating comprehension**, while you do the meaning-making.

7.4 The 4-Step Response to Objections

When someone is afraid, angry, or suspicious about AI, use this:

1. **Empathize**

    - Name the value they are protecting.
    - "You're guarding something sacred—your freedom, your faith, your livelihood. That matters."

2. **Educate**

    - Offer evidence and nuance (e.g., "The mark of the beast is about forced worship and coercion, not chips in themselves.").

3. **Bridge**

    - Show how your framework fights the same enemy.
    - "This is exactly why we want AI to remain optional and transparent—so no one can erase your conscience."

4. **Invite**

    - Ask them into partnership.
    - "What would a fair, humane version of this look like to you? Will you help shape that?"

PART 8 — MOVEMENT IMPLEMENTATION (DEPENDENCIES MADE EXPLICIT)

For the Guardians of Creation & Conscience movement (or similar efforts), success depends on specific **operational commitments**.

8.1 Minimum Inputs

- **Founder/Leader Time:** ~20–30 hours/week for the first 90 days.
- **Ambassadors:** 10–15 committed people giving 5–20 hours/month.
- **Budget:**
    - Phase 1 (0–3 months): $0–500 (tools, basic design).
    - Phase 2 (3–6 months): $1,000–2,000 (ads, occasional contractors).
    - Phase 3 (6+ months): $2,000–5,000/month (part-time coordinator, ongoing content, events).

8.2 Critical Success Factors

1. **Ambassador Training Quality**

    - They understand the three pillars (no inner life, love as action, stewardship boundaries).
    - They can use the 4-step response authentically, not mechanically.
    - They know when to escalate emotional or theological issues.

2. **Cadence and Consistency**

    - Weekly content rhythm (streams, posts, newsletters) is sustained for 13+ weeks.
    - Discord or community spaces are moderated kindly and promptly.

3. **Boundary Clarity**

    - Public stance that this framework is **not therapy**.
    - Clear guidance for users about when to seek human or professional help.

4. **Feedback and Iteration**

    - Regular debriefs with ambassadors and early adopters.
    - Willingness to adjust language and tools for different communities and cultures.

PART 9 — WEEKLY PRACTICE (SIMPLE RHYTHM)

A suggested weekly rhythm for individuals or small groups:

Daily (15–20 minutes):

- Use the **Learning Amplifier** on a relevant topic.
- Or complete one **Decision Exoskeleton** for a real decision.

2x per week (30–45 minutes):

- Run an **Emotional Debrief** for any intense AI interaction.
- Discuss one **scenario** from your playbook (e.g., environmental concerns, job fears).

Weekly (60–90 minutes):

- Group session: share wins, concerns, and updated boundaries.
- Review one Non-Negotiable and identify where your tools succeed or fail.

Over weeks, this produces:

- Better decisions, more explicitly grounded in values.
- Less anxiety and confusion.
- Stronger community capacity to discern together.

PART 10 — CLOSING: LOVE AS THE RELENTLESS CRITERION

This framework is **pro-AI insofar as AI serves love** and **anti-AI insofar as AI undermines love**.

- When AI helps heal, educate, and free people to live more fully human lives, within environmental and social limits—that is worth amplifying.
- When AI accelerates exploitation, deception, coercion, or environmental ruin—that is worth resisting.

Love, here, is not soft. It is **steely** about boundaries and consequences. It:

- Says "no" when systems demand obedience against conscience.
- Says "no" when tools exhaust creation without proportionate good.
- Says "no" when people are used as data points rather than treated as ends in themselves.

But it also:

- Says "yes" to tools that genuinely reduce suffering.
- Says "yes" to knowledge and innovation in service of justice and care.
- Says "yes" to communities learning together how to live wisely in a changing world.

The future of AI is not fixed. It will be shaped by:

- Who shows up.
- Which values they insist on.
- What trade-offs they refuse to make.

This framework is an invitation:

> **To treat AI not as a god to worship, nor a demon to fear, but as a powerful tool that must be constrained, guided, and evaluated by love.**

Use it. Adapt it. Share it. Update it as reality changes.

And above all—**keep humans, and the more-than-human world, at the center of what matters.**

END OF DOCUMENT